

H₂O.ai

<https://www.h2o.ai>

R - DOCUMENTATION



1. Overview

- Written in **JAVA**.
- Can be implemented on top of **HADOOP** (internally it runs Map/Reduce).
- Uses Java Fork/Join for **multi-threading**.
- Uses distributed Key/Value store to access and reference data.
- Gives **REST API** call for **R** and **Python**.
- **Advantages**

- **Fast**
- **Distributed**
- **In-Memory**
- **Scalable**

2. Installation

- <http://h2o-release.s3.amazonaws.com/h2o/rel-ueno/3/index.html>

The following two commands remove any previously installed H2O packages for R.

```
if ("package:h2o" %in% search()) { detach("package:h2o", unload=TRUE) }  
if ("h2o" %in% rownames(installed.packages())) { remove.packages("h2o") }
```

Next, we download packages that H2O depends on.

```
if (! ("methods" %in% rownames(installed.packages()))) { install.packages("methods") }  
if (! ("statmod" %in% rownames(installed.packages()))) { install.packages("statmod") }  
if (! ("stats" %in% rownames(installed.packages()))) { install.packages("stats") }  
if (! ("graphics" %in% rownames(installed.packages()))) { install.packages("graphics") }  
if (! ("RCurl" %in% rownames(installed.packages()))) { install.packages("RCurl") }  
if (! ("jsonlite" %in% rownames(installed.packages()))) { install.packages("jsonlite") }  
if (! ("tools" %in% rownames(installed.packages()))) { install.packages("tools") }  
if (! ("utils" %in% rownames(installed.packages()))) { install.packages("utils") }
```

Now we download, install and initialize the H2O package for R.

```
install.packages("h2o",  
                 type="source",  
                 repos=(c("http://h2o-  
release.s3.amazonaws.com/h2o/rel-ueno/3/R")))  
library(h2o)
```

3. Basic Commands

- `library(h2o)`

- load h2o in R

- `h2o.init(nthreads = -1, max_mem_size = '8G')`

- `nthreads = -1` means it uses all cores in all CPUs
- `max_mem_size` = memory to be used for H2O
- **Note:** Use all CPUs and full RAM and don't run any browsers while H2O operations is happening.

- `h2o.clusterInfo()`

- prints the cluster status and information

- `h2o.shutdown(prompt = F)`

- Shutdown the cluster
- `prompt = F` won't ask for confirmation

- `as.h2o(d)`

- converts `d` to a h2o frame (where `d` = `data.table` or `data.frame`)

- `h2o.ls()`

- to check objects in h2o environment

- `h2o.importFile("full-path-of-the-file", sep = ",")`

- similar to `read.csv`

- `h2o.exportFile(object, "full-path-for-the-file")`

- similar to `write.csv`

- `h2o.merge(x, y)`

- `h2o.rbind(x, y)`

- `h2o.cbind(x, y)`

- `h2o.table(x, y)`

- `h2o.group_by`

- `h2o.runif(100)`
- `h2o.rm("object-in-quotes")`
- `h2o.which(x == 1)`
- `h2o.ifelse`
- `h2o.strsplit`
- `h2o.tolower`
- `h2o.toupper`

- `h2o.trim`

- trims leading and trailing white spaces in a string

- `h2o.gsub`

- `h2o.impute(data, "column-a", method = mean/mode/median, by = c("column-b", "column-c"))`

- the above command imputes missing values in "column-a" with mean or median or mode grouped by "column-b", "column-c"

- `h2o.splitFrame(data, ratios = 0.5)`

- splits data into 2 parts of 0.5 each
- if ratios = c(0.75, 0.15), then splits data into 3 parts 75%, 15% and 10% (train, test and validation)

4. Modelling Commands [Explained in later sections]

- `h2o.glm`
- `h2o.gbm`
- `h2o.randomForest`
- `h2o.naiveBayes`
- `h2o.deeplearning`
- `h2o.kmeans`
- `h2o.betweenss`
- `h2o.centers`
- `h2o.predict`
- `h2o.accuracy`
- `h2o.confusionMatrix`
- `h2o.performance`
- `h2o.mse`
- `h2o.grid`

5. GLM (Generalized Linear Models)

- Extension of traditional linear models.
- Scales well for large datasets.
- Handles classification and Prediction by using family.
- Handles Gaussian, Poisson and various other family of distributions of data.
- **In case of Linear Regression:**

- $Y = X^T \beta + \beta_0 + \epsilon$

- where,
- Y = Target (independent variable) and $Y \sim N(X^T \beta + \beta_0, \sigma^2)$
- X^T = Feature vector
- β = Parameter vector
- β_0 = Intercept
- ϵ = noise, $\epsilon \sim N(0, \sigma^2)$ i.e. Gaussian random variable

- This assumes:

- Normality of error term
- Constant variance
- Additivity of covariates

- Thus, we need more flexible model (Hence, we use GLM)

- It allows variance to vary
- Allows non-linear relationship between response and covariates.

- **Aim:**

- For Classification: To maximize log-likelihood over parameter (B).

- i.e. $\max_{\beta}(\log \text{likelihood})$

- For Regression: To minimize MSE.

- GLM also uses regularization and introduces parametric regularization penalty to prevent over fitting.

- i.e. $\max_{\beta}(\log \text{likelihood} - \text{reg penalty})$

- **Parameters of regularization:**

- **alpha:** elastic net parameter [0, 1]
- **Lambda:** regularization coefficient (preferred to do *lambda_search*)

- **Model Validation:**

- Precision
- Recall
- Accuracy
- $AIC = 2k - 2\log(l(y; \mu))$

- where,

- k = no. of parameters
- $l(y; \mu)$ = fitted maximum likelihood
- It is used to compute qualities of models
- Lower the AIC value, optimal the model
- It increases penalty with increase in number of parameters, helps in preventing over fitting

- *Deviance = (max. loglikelihood.of.fitted – max.loglikelihood.of.saturated.models)*

- **Regularization:**

- Penalties of coefficients are introduced to prevent over fitting.
- Reduces the variance of prediction error.
- Handles correlated predictor variables.
- **Examples:**

- **Lasso Regression:**

- Uses L1 Regularization
- $||B||_1 = \sum_{k=1}^n |B_k|$
- if lambda value is high, all coefficients are set to 0
- Use it if the no. of features are high/correlated (Selects only one among correlated variable and sets other to 0)
- Helps in reducing sparsity

- **Ridge Regression:**

- Uses L2 Regularization
- $\|B\|_2^2 = \sum_{k=1}^n B_k^2$
- Helps in achieving numerical stability
- Shrinks all parameters proportionally
- Reduces coefficients as the penalty increases without setting any value to 0
- Reduces coefficients if the features are correlated.

- **Elastic Net:**

- Uses L1, L2 Regularization
- has alpha, lambda
- where:
 - alpha [0, 1] - elastic penalty distribution
 - lambda (> 0) - Penalty Strength

lambda	alpha	results
0	any value	no regularization
> 0	0	ridge regression
> 0	1	lasso regression
> 0	(0, 1)	elastic net

- Handles sparsity and achieves stability

- **Model Families:**

- **Gaussian (Regression)**

```
h2o.glm(x, y, training_frame = df, family = 'gaussian')
```

- **Binomial (Logistic regression)**

```
h2o.glm(x, y, training_frame = df, family = 'binomial')
```

- **Multinomial (multi-class classification)**

```
h2o.glm(x, y, training_frame = df, family = 'multinomial')
```

- **Poisson (if Y >= 0 and errors have poisson distribution)**

```
h2o.glm(x, y, training_frame = df, family = 'poisson')
```


- **Gamma (if $Y > 0$)**

```
h2o.glm(x, y, training_frame = df, family = 'gamma')
```

- **Tweedie** (includes gamma, normal, poisson, and their combination, and for $Y \geq 0$)

- **Parameters in h2o.glm:**

- training_frame
- validation_frame
- lambda
 - = **0** as default
- lambda_search
 - use **TRUE**, for searching optimal lambda value
- max_active_predictors
 - uses when lambda_search = TRUE
 - if = **10**, then stops after having 10 predictor values
- remove_collinear_columns
 - use **TRUE** to remove collinear columns
- standardize
 - use **TRUE** to scale the data to 0 mean and unit variance
- compute_p_values
 - use **TRUE** to compute p values
 - used for hypothesis testing
 - higher the p values, unreliable the features are
 - lower the significant
- nfolds
 - = **5** (does 5 fold cross validation)
 - to **print** use,

```
fit@model$training_metrics@metrics$AUC
fit@model$cross_validation_metrics@metrics$AUC
```

- **grid search over alpha:**

```
alphas = list(list(0), list(0.25), list(0.5), list(0.75),
list(1))
```

```

hyper_params = list(alpha = alphas)

grid = h2o.grid(model = 'glm', hyper_params = hyper_params,
x, y, training_frame = df, family = 'binomial')

models = lapply(grid@model_ids, function(each_model) {

  h2o.getModel(each_model)

})

print(models[[1]]@model$model_summary$regularization)

print(h2o.auc(models[[1]]))

```

- **Model Statistics:**

- h2o.mse
- h2o.rmse
- h2o.r2
- h2o.auc
- h2o.aic
- h2o.logloss
- h2o.null_deviance
 - deviance from NULL model
- h2o.residual_deviance
 - deviance of built model

- `print(fit@model$coefficients)`

- prints model coefficients

- `print(fit@model$scoring_history)`

- prints likelihood, iteration, objective scores at different timestamps and iterations

- `h2o.confusionMatrix(fit, valid = T)`

- valid = for validation data flag (TRUE or FALSE)

- `h2o.predict(fit, newdata = test)`

- used for prediction on new data set

```
o h2o.performance(fit, newdata = test)
```

- used to calculate different performance measures on the new data

6. Random Forest

- Distributed in case of H2O
- Multiple Trees (multiple weak learners)

```
h2o.randomForest(model_id          ='rf-model',          training_frame,
validation_frame, nfolds, x, y
, ntrees # grid search (10 -500)
, max_depth # grid search seq(1, 29, 2)
, min_rows # default = 10, 500 means it needs 500 TRUE and 500
FALSE to make split
# grid search seq(1, 20, 1)
, min_split_improvement # default = 10^-5
# grid search (10^-10 .... 10^-3)
, nbins # bin numeric value, grid search (8, 16, 32, 64, 128,
256, 512)
, nbins_cats # bins for categorical, grid search (8, 16, 32, 64,
128, 256, 512, 1024, 2048, 4096)
, nbins_top_level # bins to use at top level of trees
# nbins, nbins_cats, nbins_top_level less values less the
overfitting
, stopping_rounds # for specific no.of rounds if there is no
improvement, then stop
, stopping_metric # logloss, rmse, auc, mse
, stopping_tolerance # 1e-3
, seed # to reproduce
, histogram_type
#      AUTO
```

```
#    UniformAdaptive
#    Random
#    QuartilesGlobal
#    Round Robin
, sample_rate # (0, 1)
# - improves generalization
# - for large datasets use 0.7, 0.8, 1
# - for imbalanced
, sample_rate_per_class
# - c(1, 0.5) = down samples class 2 by 50%
, col_sample_rate # (0, 1)
# - improves generalization
# - reduces validation error
# - no. of cols to consider for each split
, col_sample_rate_per_tree
# if no. of cols (n) = 100
# - if col_sample_rate_per_tree = 0.75
# - and if col_sample_rate = 0.8
# - then for each tree =  $0.75 * 100 = 75$  cols are considered
# - and for each split =  $75 * 0.8 = 60$  cols are considered
, mtries # (-1 or  $\geq 1$ )
# - no. of cols to select randomly at each level
# - grid search for best value)
```

7. K Means

```
h2o.kmeans(x, k = 3, training_frame, validation, seed
, init
# Random - chooses k points randomly
# Furthest - Choose m1, calculate distance from all N-1 Points,
choose far one, repeat till K
# PlusPlus - Choose m1, calculate distance from all N-1 Points,
score by giving weights as distances,
# now pick m2 randomly out of weighted
probability, repeat till K
# User - User defined centers
, estimate_k # flag that helps and builds till <= K
, max_iterations # 1 - 10^6
, standardize = TRUE)
```

- Determine K
 - use estimate_k parameter
 - internally it calculates:

- $PRE = (SSW[\text{after split}] - SSW[\text{before split}]) / SSW[\text{before split}]$
 - if $PRE < \text{threshold}$, stop

- where, $\text{threshold} = \min(0.8, 0.02 + (10/\# \text{ of rows}) + (2.5/(\# \text{ of features})^2))$

8. PCA

- transforms correlated features to uncorrelated features (principle components)
- maximizes variances, reduces covariance
- used for dimension reduction
- makes all dimensions as orthogonal
- used before KMeans for best results
- Categorical - internally it converts to one hot encoding

```

h2o.pcomp(model_id, training_frame

          , transform

          # none

          # standardize

          # normalize

          # Demean

          # Descale

          , seed)

```

9. GLRM

- Generalized Low Rank Models
- Dimensional Reduction
- Useful to reconstruct missing values
- Find important features
- Parallel and optimized
- same use as SVD, Matrix factorization

$$\circ \quad A [m \times n] = X [m \times k] \%*\% Y [k \times n]$$

- Choose k = lower the better compressed

```
h2o.glm(training_frame, cols = 1:ncol(data), k = 10)
```

10. Load and Save Models

```

h2o.loadModel(path)

h2o.saveModel(model, dir, name)

```

11. GBM

- Gradient Boosting Model
- Gradually improved estimations
- Models Non-Linear relationships
- Weak classification algorithms are sequentially applied to the incrementally changed data to create a series of decision trees, producing an ensemble
- In H2O
 - Distributed and Parallelized
 - Fast and memory efficient
 - Uses stochastic gradient descent with column and row sampling
- Boosting is optimized by Gradient descent to minimize a model loss

- **Steps**

- a) fit a model to the data $F1(x) = Y$
- b) fit a model to the residuals $(Y - Y')$

$$h1(x) = Y - F1(x)$$

- c) create a model $F2(x) = F1(x) + h1(x)$

i. e.

$$F(x) = F1(x) \rightarrow F2(x) \text{ --- } \rightarrow FM(x)$$

$$\text{where } FM(x) = FM-1(x) + hM-1(x)$$

and $h1(x)$ weak learner (a decision stump)

- **Main Parameters**

- $\text{eta} = 0.3 \text{ or } 0.1$
- $\text{max_depth} = 3$
- $\text{n_estimators} = 100$
- $\text{subsample} = 1$

- **Tips by Owen Zhang (Kaggle #2) and Tianqi Chen**

- Tune learning rate by using fixed values of trees
- Use random search over grid search if search space is too large

GBDT Hyper Parameter Tuning

Hyper Parameter	Tuning Approach	Range	Note
# of Trees	Fixed value	100-1000	Depending on datasize
Learning Rate	Fixed => Fine Tune	[2 - 10] / # of Trees	Depending on # trees
Row Sampling	Grid Search	[.5, .75, 1.0]	
Column Sampling	Grid Search	[.4, .6, .8, 1.0]	
Min Leaf Weight	Fixed => Fine Tune	3/(% of rare events)	Rule of thumb
Max Tree Depth	Grid Search	[4, 6, 8, 10]	
Min Split Gain	Fixed	0	Keep it 0

Best GBDT implementation today: <https://github.com/tqchen/xgboost>
by **Tianqi Chen** (U of Washington)



- **Codes**

```
fit = h2o.gbm(x, y, training_frame, validation_frame, nfolds
, ntrees # default = 50
, max_depth # default = 5
, min_rows # default = 10
, min_split_improvement
, nbins
, nbin_cats
, seed
, learn_rate # (0, 1)
# - lower the better but more time to reach minima
# - use 0.05, with 0.99 learn_rate_annealing
, learn_rate_annealing
```



```
# - reduces learning rate by this factor for every tree

# - start with learn_rate

# - ends with learn_rate * (learn_rate_annealing) ^ N

# - instead of LR = 0.01, use LR = 0.05, LRA = 0.99 -
executes faster

# - use = 1 to disable, but we should use low learn_rate

, distribution # loss function

# - AUTO - if selected and if y == numeric, it selects
gaussian, else bernoulli

# - bernoulli - 2 class

# - multinomial - multi class

# - gaussian - numeric

# - poisson

# - gamma

# - laplace

# - quantile

# - huber

# - tweedie

, sample_rate

, sample_rate_per_class

, col_sample_rate

, col_sample_rate_per_tree
```

```
, score_each_iterator # Flag

, score_tree_interval # if = 5, scores every 5 trees

, weights_column # no. of times each row to be repeated

, balance_classes # flag

, class_sampling_factors # c(0.5, 1) = reduces class - 1
by 50%

, max_after_balance_size # if 0.85 reduces the total data
size to 85%

# if 1.7 then it increases the total data size to 117%

, stopping_rounds

, stopping_metric

, stopping_tolerance

, max_runtime_secs

, checkpoint

# used to continue model building

)

h2o.varimp(fit)
```

12. Deep Learning

- Neural Networks that are having more hidden layers
- Fast, memory efficient
- Multi-threaded, parallel and distributed
- Implemented in Java
- **Adaptive Learning rate**
 - if $\text{error}_N > \text{error}_{N-1}$
 - $\alpha = \alpha/2$
- Regularization - L1 ($\lambda_1 * \text{Sigma}(W)$), L2 ($\lambda_2 * \text{Sigma}(W^2)$), dropout, model averaging, HOGWILD!
- Can have learning rate, momentum
- Auto encoders for unsupervised
- **Activation Functions**
 - tanh: $(\exp(x) - \exp(-x))/(\exp(x) + \exp(-x))$
 - makes the value of x in range of [-1, 1]
 - relu: $\max(0, x)$
 - rectified linear unit
 - maxout: $\max(x_1, x_2)$
- **Loss functions:**
 - for gaussian - use mse
 - for laplace - use absolute error
 - for bernoulli - use cross entropy
- **Sample:**

```
h2o.deeplearning(x, y, training_frame, validation_frame,  
distribution = 'bernoulli')
```

- **Loss minimization:**
 - Uses SGD to parallelize via back propagation
 - It also uses HOGWILD! to parallelize as SGD can't be parallelized completely
 - HOGWILD! is lock free parallelization scheme

- HOGWILD! uses shared memory when multiple cores handle multiple subsets of data
- **Other Optimizations:**
 - L1
 - L2
 - dropout
 - input_dropout_ratios
 - hidden_dropout_ratios
 - momentum (0, 1)
 - helps in avoiding local minima
 - too much momentum leads instability
 - momentum_start
 - momentum_ramp
 - momentum_stable
 - **nesterov_accelerated_gradient** # recommended
 - Large momentum + low learning rate is better
 - damps oscillation in the direction of high curvature
 - Rate Annealing
 - 10^{-6}
 - i.e. it takes 10^6 samples to change learning rate to half
 - Adaptive Learning
 - avoids slow convergence
 - adds benefits of momentum + rate_annealing
 - rho -> momentum (0.9 or 0.9999)
 - epsilon -> rate annealing (10^{-10} or 10^{-4})
- **Code:**

```
h2o.deeplearning(x, y, training_frame, validation_frame, seed,
nfolds

, activation

# Tanh

# Tanhwithdropout
```

```
# Rectifier

# Rectifierwithdropout

# maxout

# maxoutwithdropout

, hidden # c(100, 100) - 2 hidden layers with 100 neurons each

, epochs # no. of times to iterate through the data

, variable_importance # flag

, weight_columns

, balance_classes

, class_sampling_factors

, max_after_balance_size

, standardize

, checkpoint

, adaptive_rate # flag

, input_dropout_ratio # (0.1 or 0.2 suggested)

, hidden_dropout_ratios # [0, 1) - default = 0.5

, l1

, l2

, loss

# - Automatic

# - CrossEntropy
```

```
# - Quadratic

# - Huber

# - Absolute

, distribution

, score_interval

, stopping_rounds

, stopping_metric

, stopping_tolerance

, autoencoder # flag

, rho # adaptive rate = T, then use this - adaptive rate time
decay factor

, epsilon # adaptive rate time smoothing factor

, shuffle_training_data

, missing_values_handling # Skip/MeanImputation

, rate

, rate_annealing

, rate_decay

, nesterov_accelerated_gradient # flag

, momentum_start # 0.5 is preferred

, momentum_ramp

, momentum_stable)
```

- **Tuning:**
 - Use hidden:
 - [200, 200]
 - [512]
 - [64, 64, 64]
 - [32, 32, 32, 32, 32]
 - Use L1, L2, dropout, adaptive rate with rho, epsilon

13. Stacked Ensembles

- **Steps:**
 - Specify a list of L base algorithms with model parameters
 - Specify a meta learning algorithm
 - Train L models on training
 - Predict on validation set (N rows)
 - we get N x L matrix (level one data)
 - Use the level one data to train a super meta learner
 - Predict on test data
- **Code:**

```
fit = h2o.stackedEnsemble(x, y, training_frame, model_id,
base_models = list(fit1@model_id, fit2@model_id))

fit = h2o.stackedEnsemble(x, y, training_frame, model_id,
base_models = list(gbm_grid_fit@model_ids))
```

- **Using h2oEnsemble**

```
library(h2oEnsemble)

learner = c('h2o.randomForest', 'h2o.deeplearning.1',
'h2o.deeplearning.2')

metalearner = 'h2o.glm.wrapper'
```

```
fit = h2o.ensemble(x, y, training_frame, family = 'binomial',  
learner = learner,  
  
metalearner = metalearner, cvControl = list(V = 5))  
  
h2o.ensemble_performance(fit, newdata)  
  
predict(fit, newdata)
```

- **Default Wrappers:**

- h2o.glm.wrapper
- h2o.gbm.wrapper
- h2o.randomForest.wrapper
- h2o.deeplearning.wrapper

14. H2O Parameter Tuning

- `hyper_parameters = list(...)`

- **Grid Search**

```
grid = h2o.grid('gbm', x, y, training_frame, hyper_params =
hyper_parameters)

model = lapply(grid@model_ids, function(x) {

    h2o.getModel(x)

})
```

- **Random Search**

```
search_criteria = list(strategy = 'RandomDiscrete', max_runtime_sec =
600, max_models = 100, stopping_metric = 'AUTO', stopping_tolerance =
0.0001, stopping_rounds = 5, seed = 1234)

grid = h2o.grid(grid_id = 'random_search', search_criteria =
search_criteria, x, y, training_frame, 'gbm', hyper_params =
hyper_parameters)

h2o.getGrid(grid_id = 'random_search', sort_by = 'rmse')
```